

# HPC tools for programming

Victor Eijkhout

2020



# Justification

High Performance Computing requires, beyond simple use of a programming language, a number programming tools. These tutorials will introduce you to some of the more important ones.

# Intro to file types

# File types

Text files	
Source Header	Program text that you write also written by you, but not really program text.
Binary files	
Object file	The compiled result of a single source file
Library	Multiple object files bundled together
Executable	Binary file that can be invoked as a command
Data files	Written and read by a program

# Text files

- Source files and headers
- You write them: *make sure you master an editor*
- The computer has no idea what these mean.
- They get compiled into programs.

(Also 'just text' files: READMEs and such)

# Binary files

- Programs. (Also: object and library files.)
- Produced by a compiler.
- Unreadable by you; executable by the computer.

Also binary data files; usually specific to a program.  
(Why don't programs write out their data in readable form?)

# Compilation



# Compilers

Compilers: a major CS success story.

- The first Fortran compiler (Backus, IBM, 1954): multiple man-years.
- These days: semester project for graduate students.  
Many tools available (`lex`, `yacc`, `clang-tidy`)  
Standard textbooks ('Dragon book')
- Compilers are very clever!  
You can be a little more clever in assembly – maybe  
but compiled languages are  $10\times$  more productive.

## Compilation vs interpreted

- Interpreted languages: lines of code are compiled 'just-in-time'.  
Very flexible, sometimes very slow.
- Compiled languages: code is compiled to machine language:  
less flexible, very fast execution.
- Virtual machine: languages get compiled to an intermediate language  
(Pascal, Python, Java)  
pro: portable; con: does not play nice with other languages.
- Scientific computing languages:
  - Fortran: pretty elegant, great at array manipulation  
Note: Fortran2003 is modern; F77 and F90 are not so great.
  - C: low level, allows great control, tricky to use
  - C++: allows much control, more protection, more tools  
(kinda sucks at arrays)

# Simple compilation

hello.c

```
int main() {  
    printf("Hello world\n");  
    return 0;  
}
```

icc -o hello.exe hello.c

hello.exe

- From source straight to program.
- Use this only for short programs.


```
%% gcc hello.c  
%% ./a.out  
hello world
```

```
%% gcc -o helloprog hello.c  
%% ./helloprog  
hello world
```

# Exercise 1

```
/******  
****  
**** This program source is part of  
**** Introduction to High-performance Scientific Computing  
**** by Victor Eijkhout  
**** copyright Victor Eijkhout 2011-2020  
****  
**** hello.c : simple hello world program  
****  
*****  
#include <stdlib.h>  
#include <stdio.h>  
  
int main() {  
    printf("hello world\n");  
    return 0;  
}
```

# Separate compilation



- Large programs best broken into small files,
- ... and compiled separately (can you guess why?)
- Then 'linked' into a program; linker is usually the same as the compiler.

# Exercise 2

Main program: fooprogram.c

Subprogram: foosub.c

```
/*
```

```
*****
```

```
/*
```

```
*****
```

```
****
```

```
**** This program source is
part of
```

```
**** Introduction to High-
performance Scientific
Computing
```

```
**** by Victor Eijkhout
```

```
**** copyright Victor
Eijkhout 2011-2020
```

```
****
```

```
**** fooprogram.c : simple main
program using external
function
```

```
****
```

```
****
```

```
**** This program source is
part of
```

```
**** Introduction to High-
performance Scientific
Computing
```

```
**** by Victor Eijkhout
```

```
**** copyright Victor
Eijkhout 2011-2020
```

```
****
```

```
**** foosub.c : declaration
of function for fooprogram.c
```

```
****
```

```
*****
```

```
*/
```

```
*/
```

# Compiler options 101

- You have just seen two compiler options.
- Commandlines look like

```
command [ options ] [ argument ]
```

where square brackets mean: 'optional'

- Some options have an argument

```
icc -o myprogram mysource.c
```

- Some options do not.

```
icc -g -o myprogram mysource.c
```

- Question: does `-c` have an argument? How can you find out?

```
icc -g -c mysource.c
```

# Object files

- Object files are unreadable. (Try it. How do you normally view files? Which tool sort of works?)
- But you can get some information about them.

```
#include <stdlib.h>
#include <stdio.h>
void bar(char *s) {
    printf("%s", s);
}
```

```
[c:264] nm foosub.o
0000000000000000 T _bar
U _printf
```

Where T: stuff defined in this file  
U: stuff used in this file



## Compiler options 102

- Optimization level: -O0, -O1, -O2, -O3  
(‘I compiled my program with oh-two’)  
Higher levels usually give faster code. Level 3 can be unsafe.  
(Why?)
- 
- -g is needed to run your code in a debugger. Always include this.
- The ultimate source is the ‘man page’ for your compiler.

# Example of optimization

Givens program

```
// rotate.c
void rotate(double *x,double *y,double alpha) {
    double x0 = *x, y0 = *y;
    *x = cos(alpha) * x0 - sin(alpha) * y0;
    *y = sin(alpha) * x0 + cos(alpha) * y0;
    return;
}
```

Run with optimization level 0,1,2,3 we get:

Done after 8.649492e-02

Done after 2.650118e-02

Done after 5.869865e-04


Done after 6.787777e-04

## Exercise 3

The file `rotate.c` can be speeded up by compiler transformations.  
Compile this file with optimization levels `0, 1, 2, 3`  
(try both the Intel and gcc compilers)  
observe run time and conjecture what transformations can explain this.  
  
Write a report of your investigations.

# Libraries

# Libraries



- Sometimes you have many object files: convenient to bundle them
- Easier to link to
- Easy to distribute as a product.
- Software library: collection of object files that can be linked to a main program.

## Static / non-shared libraries

- Static libraries are created with `ar`
- Inspect them with `nm`
- Link as object file:

```
icc -o myprogram main.o ../lib/libfoo.a
```

- Or:

```
icc -o myprogram main.o -L../lib -lfoo.a
```

```
mkdir ../lib                                %% nm ../lib/libfoo.a
ar cr ../lib/libfoo.a foosub.o              ../lib/libfoo.a(foosub.o):
                                             00000000 T _bar
                                             U _printf
```

## Static library example

Use `ar` to add object files to `.a` file.

```
icc -g -O2 -std=c99 -c foosub.c
for o in foosub.o ; do \
    ar cr libs/libfoo.a ${o} ; \
done
icc -o staticprogram fooprogram.o -Llibs -lfoo
-rwx----- 1 eijkhout G-25072 38192 Sep 23 18:15 staticprogram
./staticprogram
hello world
```

# Dynamic/shared libraries

Created with the compiler,  
-shared flag.

```
icc -O2 -std=c99 -fPIC -c foosub.c  
icc -o libs/libfoo.so -shared foosub.o  
icc -o dynamicprogram fooprogram.o -Llibs -lfoo
```



## Executable size

Static libraries are baked into the executable  
shared libraries are linked at runtime.

```
# Making static library
icc -o staticprogram fooprogram.o -Llibs -lfoo
-rwx----- 1 eijkhout G-25072 28232 Sep 23 14:25 staticprogram
# Using dynamic library
icc -o dynamicprogram fooprogram.o -Llibs -lfoo
-rwx----- 1 eijkhout G-25072 28160 Sep 23 14:25 dynamicprogram
```

## Needs something more

Program can not immediately be run.

Use `ldd` to see what libraries it needs:

```
./dynamicprogram: error while loading shared libraries:  
libfoo.so: cannot open shared object file: No such file or directory
```

```
ldd dynamicprogram | grep libfoo  
libfoo.so => not found
```

## The ell-dee library path

Libraries are found by updating the `LD_LIBRARY_PATH`:

```
export LD_LIBRARY_PATH=${LD_LIBRARY_PATH}:/libs
ldd dynamicprogram | grep libfoo
    libfoo.so => ./libs/libfoo.so (0x00002ad6604c1000)
./libs dynamicprogram
hello world
```

# The rpath

You can also bake the path into the program:

```
icc -O2 -std=c99 -fPIC -c foosub.c
icc -o libs/libfoo.so -shared foosub.o
icc -o rpathprogram fooprogram.o \
    -Wl,-rpath=./libs -Llibs -lfoo
-rwx----- 1 eijkhout G-25072 28160 Sep 23 13:41 rpathprogram
./rpathprogram
hello world
```

(Notice the bizarre combination of minuses and commas)

# Cross-language linking



# **Profiling and debugging; optimization and programming strategies.**

# Analysis basics

- Measurements: repeated and controlled  
beware of transients, do you know where your data is?
- Document everything
- Script everything



# Compiler options

- Defaults are a starting point
- use reporting options: `-opt-report`, `-vec-report`  
useful to check if optimization happened / could not happen
- test numerical correctness before/after optimization change  
(there are options for numerical correctness)

# Optimization basics

- Use libraries when possible: don't reinvent the wheel
- Premature optimization is the root of all evil (Knuth)

# Code design for performance

- Keep inner loops simple: no conditionals, function calls, casts
- Avoid small functions: try macros or inlining
- Keep in mind all the cache, TLB, SIMD stuff from before
- SIMD: Fortran array syntax helps

## Multicore / multithread

- Use `numactl`: prevent process migration
- 'first touch' policy: allocate data where it will be used
- Scaling behaviour mostly influenced by bandwidth

# Multinode performance

- Influenced by load balancing
- Use HPCtoolkit, Scalasca, TAU for plotting
- Explore 'eager' limit (mvapich2: environment variables)

# Classes of programming errors

Logic errors:

functions behave differently from how you thought,  
or interact in ways you didn't envision

Hard to debug

Coding errors:

send without receive

forget to allocate buffer

Debuggers can help

## Defensive programming



# Defensive programming

- Keep It Simple ('restrict expressivity')
- Example: use collective instead of spelling it out
- easier to write / harder to get wrong  
the library and runtime are likely to be better at optimizing than you

# Memory management

Beware of memory leaks:

keep allocation and free in same lexical scope

# Modular design

Design for debuggability, also easier to optimize

Separation of concerns: try to keep code aspects separate

Premature optimization is the root of all evil (Knuth)

# MPI performance design

Be aware of latencies: bundle messages  
(this may go again separation of concerns)

Consider 'eager limit'

Process placement, reduction in number of processes

## Debugging

# Debugging

I assume you know about gdb and valgrind. . .

- Interactive use of gdb, starting up multiple xterms feasible on small scale
- Use gdb to inspect dump:  
can be useful, often a program crashes hard and leaves no dump

Note: compile options `-g -O0`

# Parallel debuggers

Defensive programming  
Debugging

Allinea DDT 4.2-34404

File View Control Search Tools Window Help

Current Group: All Focus on current: Group Process Thread Step Threads Together

0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15

Create Group

Project Files

Search (Ctrl+K)

Application Code

Sources

problem1.c

External Code

```
18 MPI_Finalize();
19 MPI_Barrier(comm);
20 }
21 }
22
23 int main(int argc, char **argv) {
24     MPI_Comm comm;
25
26     MPI_Init(&argc, &argv);
27     comm = MPI_COMM_WORLD;
28
29     loop_for_await(comm);
30
31     MPI_Finalize();
32     return 0;
33 }
34
```

Locals Current Line(s) Current Stack

Current Line(s)

Variable Name	Value
argc	1
argv	0x7ffffff7958

Type: none selected

Input/Output Breakpoints Watchpoints Stacks Tracepoints Tracepoint Output Logbook Evaluate

Stacks

Processes	Function
16	main (problem1.c:26)

Expression Value

Ready

## Buggy code

```
for (it=0; ; it++) {  
    double randomnumber = ntids * ( rand() / (double)RAND_MAX  
    printf("[%d] iteration %d, random %e\n",mytid,it,randomnum  
    if (randomnumber>mytid && randomnumber<mytid+1./(ntids+1))  
        MPI_Finalize();  
    MPI_Barrier(comm);  
}
```



# Parallel inspection

File View Control Search Tools Window Help

Current Group: All Focus on current: Group Process Thread Step Threads Together

Create Group

Project Files

Search (Ctrl+K)

Application Code

Sources

problem1.c

loop\_for\_await

main(int argc)

External Code

problem1.c

```

10 // Initialize the random number generator
11 srand((int)(mytid*(double)RAND_MAX/ntids));
12
13
14 for (it=0; ; it++) {
15     double randomnumber = ntids * ( rand() / (double)RAND_MAX );
16     printf("[%d] iteration %d, random %e\n", mytid, it, randomnumber);
17     if (randomnumber > mytid && randomnumber < mytid+1./ntids)
18         MPI_Finalize();
19     MPI_Barrier(comm);
20 }
21
22
23 int main(int argc, char **argv) {
24     MPI_Comm comm;
25
26     MPI_Init(&argc, &argv);
27     comm = MPI_COMM_WORLD;

```

Locals

Variable Name	Value
comm	1140850688
it	31
mytid	1
ntids	16
randomnumber	1.056087621

Stacks

Processes	Function
16	main (problem1.c:29)
1	loop_for_await (problem1.c:18)
15	loop_for_await (problem1.c:19)
15	MPI_Barrier (barrier.c:411)
15	MPIR_Barrier_impl (barrier.c:266)
15	MPIR_Barrier_MV2 (barrier_osu.c:198)
15	MPIR_Barrier_intra_MV2 (barrier_osu.c:166)
1	MPIR_shmem_barrier_MV2 (barrier_osu.c:104)
1	MPIDI_CH3I_SHMEM_COLL_Barrier_gather (ch3_shmem_coll.c:940)
1	MPIDI_CH3I_Progress_test (ch3_progress.c:471)
1	MPIDI_CH3I_SMP_read_progress (ch3_smp_progress.c:743)
1	MPIDI_CH3I_SMP_pull_header (ch3_smp_progress.c:4345)
14	MPIR_shmem_barrier_MV2 (barrier_osu.c:113)

Input/Output\* Breakpoints Watchpoints Stacks Tracepoints Tracepoint Output Logbook Evaluate








Expression	Value

Ready

# Stack trace

Stacks	
Processes	Function
16	main (problem1.c:29)
1	loop_for_await (problem1.c:18)
15	loop_for_await (problem1.c:19)
15	PMPI_Barrier (barrier.c:411)
15	MPIR_Barrier_impl (barrier.c:266)
15	MPIR_Barrier_MV2 (barrier_osu.c:198)
15	MPIR_Barrier_intra_MV2 (barrier_osu.c:166)
1	MPIR_shmem_barrier_MV2 (barrier_osu.c:104)
1	MPIR_CH3I_SHMEM_COLL_Barrier_gather (ch3_shmem_coll.c:940)
1	MPIR_CH3I_Progress_test (ch3_progress.c:471)
1	MPIR_CH3I_SMP_read_progress (ch3_smp_progress.c:743)
1	MPIR_CH3I_SMP_pull_header (ch3_smp_progress.c:4345)
14	MPIR_shmem_barrier_MV2 (barrier_osu.c:113)

## Variable inspection

Locals	Current Line(s)	Current Stack
Locals  		
Variable Name	Value	
... comm	 1140850688	
... it	 31	
... mytid	 1	
... ntids	 16	
... randomnumber	 1.056087621	